

# Scalable Data Analysis with Spark and Microsoft R Server

[Ali Zaidi, alizaidi@microsoft.com](mailto:alizaidi@microsoft.com)

October 25th, 2016

# Scalable Data Analysis with R and Spark

# Unpacking the Title

- Scalable - should be able to accommodate large datasets
  - $p > n$
  - streaming datasets
  - generally large datasets where you want to do some inference
- R and Spark
  - R is great at statistical modeling, visualization and inference
  - R is also lazy, and functional
  - Let's reuse the API



# Background

- Studied statistics and machine learning at the University of Toronto
- Joined Microsoft through the Revolution Analytics acquisition last year
- Have used R for about 7 years now
- Hadoop for 2.5 years
- Spark for 1.5 years



# Strengths of R

## Where R Succeeds

- Expressive
- Open source
- Extendable – nearly 10,000 packages with functions to use, and that list continues to grow
- Focused on statistics and machine learning – utilized by academics and practitioners
- Advanced data structures and graphical capabilities
- Large user community, academics and industry
- It is designed by statisticians



# Weaknesses of R

## Where R Falls Short

- It is designed by statisticians
- Inefficient at element-by-element computations
- May make large demands on system resources, namely memory
- Data capacity limited by memory
- Single-threaded



# Microsoft R Server - Scalable R

## WODA - Reusable API

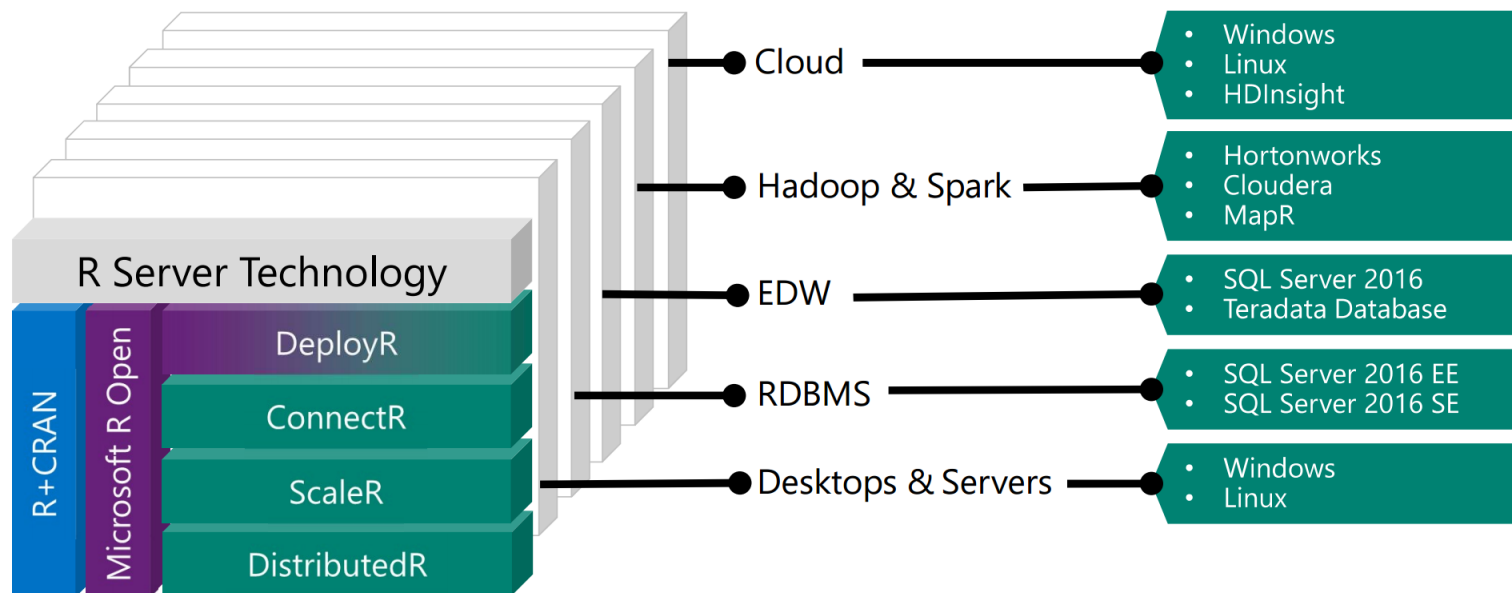
- Breaks R's memory shackles by using external data frames (XDFs) and bridges to distributed systems (Spark, HadoopMR, Teradata)
- One of the core principles behind Microsoft R Server is the WODA framework
  - WODA – write once deploy anywhere
  - Encourages API reuse
- Provides seamless transition from a local environment to a cluster environment
- Don't want to rewrite the entire code-base in Java/Scala
- Deployment should be possible within seconds



# Microsoft R Server

## Components

### Microsoft R Server





# R APIs for Spark - A Tale of Three APIs

# The aRt of Being Lazy

## Lazy Evaluation in R



- R, like its inspiration, Scheme, is a *functional* programming language
- R evaluates lazily, delaying evaluation until necessary, which can make it very flexible
- R has a high memory footprint, and can easily lead to crashes if you aren't careful

# R APIs for Spark

## RxSpark

- Allows the distribution of local R code into spark applications
- The only abstractions are those of the **RevoScaleR** package
  - data resides as distributed on-disk objects (**xdf**'s)
  - Functions and data transformations are provided by traditional R objects
- Each **RevoScaleR** function invokes it's own Spark applications and converts XDFDs into RDDs/DataFrames
  - application persists for the duration of the job to avoid JVM creation overhead
  - for multi-iteration jobs, data is cached
  - no need for the developer to write any Spark code
- Available through Azure HDInsight Clusters
  - currently utilizing Spark 1.6
- User defined functions can be applied at the data partition level using **rxExec**



# R APIs for Spark

## SparkR

- The standard R API for Spark since 1.4 (MLlib support started in 1.5)
- R package provides functions that invoke functions directly on the JVM
  - Uses a RPC server and provides JVM wrappers
- `DataFrame` support is inspired by the `dplyr` package
  - tries to emulate `dplyr` syntax, but doesn't use NSE
  - `dplyr: taxi %>% group_by(pickup_nhood) %>% summarise(ave_delay = mean(pickup_delay))`
  - `SparkR: taxi %>% group_by(taxi$pickup_nhood) %>% summarise(ave_delay = mean(taxi$pickup_delay))`
- Limited ML: `glm`, `naive-bayes`, and `kmeans`
- Spark 2.0: support for custom UDFs using `dapply` and `gapply`
  - each partition must fit into an R process on the worker node



# R APIs for Spark

Masking with `dplyr`



# R APIs for Spark

## sparklyr

- Turns out `dplyr` already has a SQL backend
- Since all Spark DataFrame operations are conducted at the Spark SQL level, utilize `dplyr`'s SQL backend rather than JVM wrappers
- 100% support for `dplyr` SQL inside of Spark, including NSE:
  - `_lyr:taxi %>% group_by(pickup_nhood) %>% summarise(ave_delay = mean(pickup_delay))`
- Full support for all SparkML
- No UDF support yet



# R APIs for Spark

`sparklyr` relies on RPC layer provided by `sparkapi`

- `sparklyr` is meant to be the DSL, providing easy data manipulation with `dplyr` and ML analogously to `stats` and other modeling packages
- The core RPC layer is not inside of `sparklyr`
- `sparkapi` provides the core R to Java RPC bridge publicly, and provides a simple extension mechanism to call arbitrary Spark APIs packages
  - e.g., extensions to connect o [H2O Sparkling Water](#) with `sparkapi` now exist



# Azure HDInsight

## Full Managed Hadoop/Spark on the Cloud

- With Azure HDInsight, you don't have to choose
- Can use any combination of the APIs for your data science application
- Focus less on optimizing code, rewriting your functions, and focus more on developing applications





# Data Manipulation Examples

# sparklyr

## Tidy Data Manipulation Using `dplyr` syntax

- For data scientists comfortable with R, using `sparklyr` requires no prior Spark knowledge
- `dplyr` statements are converted to SQL statements, sent to Catalyst and optimized



# Import Into Spark DataFrames

```
origins <- file.path("wasb://mrs-spark@alizaidi.blob.core.windows.net",  
                    "user/RevoShare/alizaidi/Freddie/Acquisition")  
  
library(sparklyr)  
sc <- spark_connect("yarn-client")  
freddie_origins <- spark_read_csv(sc,  
                                  path = origins,  
                                  name = 'freddie_origins',  
                                  header = FALSE,  
                                  delimiter = "|"  
                                  )
```



# Using dplyr

- Now it's a Spark DataFrame
- It is also of class `tbl_sql`, so all `dplyr` methods are converted to `Spark SQL` statements and run on the spark application defined through the spark context `sc`

```
class(freddie_origins)
```

```
## [1] "tbl_spark" "tbl_sql" "tbl_lazy" "tbl"
```

```
library(dplyr)
freddie_origins %>% head
```

```
## Source:   query [?? x 25]
## Database: spark connection master=yarn-client app=sparklyr local=FALSE
##
##      V1      V2      V3      V4      V5      V6      V7      V8      V9      V10     V11
##   <chr> <int> <chr> <int> <int> <chr> <int> <chr> <dbl> <chr> <int>
## 1   751 199910      N 202909      NA    000      1      0     71     20 180000
## 2   733 199909      N 202908 29540    000      1      0     51          116000
## 3   755 199905      N 202904 29540     30      1      0     95     38 138000
## 4   669 200206      N 202901      NA    000      1      0     80     33 162000
## 5   732 199904      N 202903 17140    000      1      0     25     10  53000
## 6   715 199904      N 202903 17140    000      1      0     67     35  91000
## # ... with 14 more variables: V12 <int>, V13 <dbl>, V14 <chr>, V15 <chr>,
## #   V16 <chr>, V17 <chr>, V18 <chr>, V19 <int>, V20 <chr>, V21 <chr>,
## #   V22 <int>, V23 <int>, V24 <chr>, V25 <chr>
```



# Renaming Columns

```
freddie_rename <- freddie_origins %>% rename(  
  credit_score = V1,  
  first_payment = V2,  
  first_home = V3,  
  maturity = V4,  
  msa = V5,  
  mi_perc = V6,  
  num_units = V7,  
  occ_status = V8,  
  cltv = V9,  
  dti = V10,  
  upb = V11,  
  ltv = V12,  
  orig_rate = V13,  
  channel = V14,  
  ppm = V15,  
  prod_type = V16,  
  state = V17,  
  prop_type = V18,  
  post_code = V19,  
  loan_number = V20,  
  loan_purpose = V21,  
  orig_term = V22,  
  num_borrowers = V23,  
  seller = V24,  
  servicer = V25  
)  
  
freddie_rename %>% head
```

```
## Source:   query [?? x 25]
```

```
## Database: spark connection master=yarn-client app=sparklyr local=FALSE
```



```
## credit_score first_payment first_home maturity msa mi_perc num_units
```

# Create Date Fields

The origination date is buried inside the loan number field. We will pick it out by indexing the loan number substring:

```
freddie_rename %>% select(loan_number)
```

```
## Source:   query [?? x 1]
## Database: spark connection master=yarn-client app=sparklyr local=FALSE
##
##      loan_number
##      <chr>
## 1 F199Q1000001
## 2 F199Q1000002
## 3 F199Q1000003
## 4 F199Q1000004
## 5 F199Q1000005
## 6 F199Q1000006
## 7 F199Q1000007
## 8 F199Q1000008
## 9 F199Q1000009
## 10 F199Q1000010
## # ... with more rows
```



# Substring Operations

```
freddie_rename <- freddie_rename %>%
  mutate(orig_date = substr(loan_number, 3, 4),
         year = as.numeric(substr(loan_number, 3, 2)))

freddie <- freddie_rename %>%
  mutate(orig_year = paste0(ifelse(year < 10, "200",
                                   ifelse(year > 16, "19",
                                           "20")), year))

freddie <- freddie %>%
  mutate(orig_year = substr(orig_year, 1, 4))

freddie %>% head
```

```
## Source:   query [?? x 28]
## Database: spark connection master=yarn-client app=sparklyr local=FALSE
##
##   credit_score first_payment first_home maturity   msa mi_perc num_units
##   <chr>         <int>         <chr>      <int> <int>   <chr>      <int>
## 1          751      199910          N    202909   NA       000         1
## 2          733      199909          N    202908 29540     000         1
## 3          755      199905          N    202904 29540     30         1
## 4          669      200206          N    202901   NA       000         1
## 5          732      199904          N    202903 17140     000         1
## 6          715      199904          N    202903 17140     000         1
## # ... with 21 more variables: occ_status <chr>, cltv <dbl>, dti <chr>,
## #   upb <int>, ltv <int>, orig_rate <dbl>, channel <chr>, ppm <chr>,
## #   prod_type <chr>, state <chr>, prop_type <chr>, post_code <int>,
## #   loan_number <chr>, loan_purpose <chr>, orig_term <int>,
## #   num_borrowers <int>, seller <chr>, servicer <chr>, orig_date <chr>,
## #   year <dbl>, orig_year <chr>
```



# Calculate Average Credit Score by Year

```
fico_year <- freddie %>% group_by(orig_year, state) %>%  
  summarise(ave_fico = mean(credit_score)) %>% collect  
fico_year %>% head
```

```
## Source: local data frame [6 x 3]  
## Groups: orig_year [6]  
##  
##   orig_year state ave_fico  
##   <chr> <chr>    <dbl>  
## 1    2008    NM  730.7395  
## 2    2012    WI  769.2008  
## 3    2015    MI  752.5942  
## 4    2014    AR  750.9074  
## 5    2011    KS  763.1039  
## 6    2006    WI  728.6378
```





# Summarize In a Function

```
year_state_sum <- function(val = "credit_score") {  
  
  library(lazyeval)  
  
  year_state <- freddie %>% group_by(orig_year, state) %>%  
    summarise_(sum_val = interp(~mean(var), var = as.name(val)))  
  
  year_state <- year_state %>% collect  
  
  names(year_state)[3] <- paste0("ave_", val)  
  
  return(year_state)  
  
}
```



# Plot

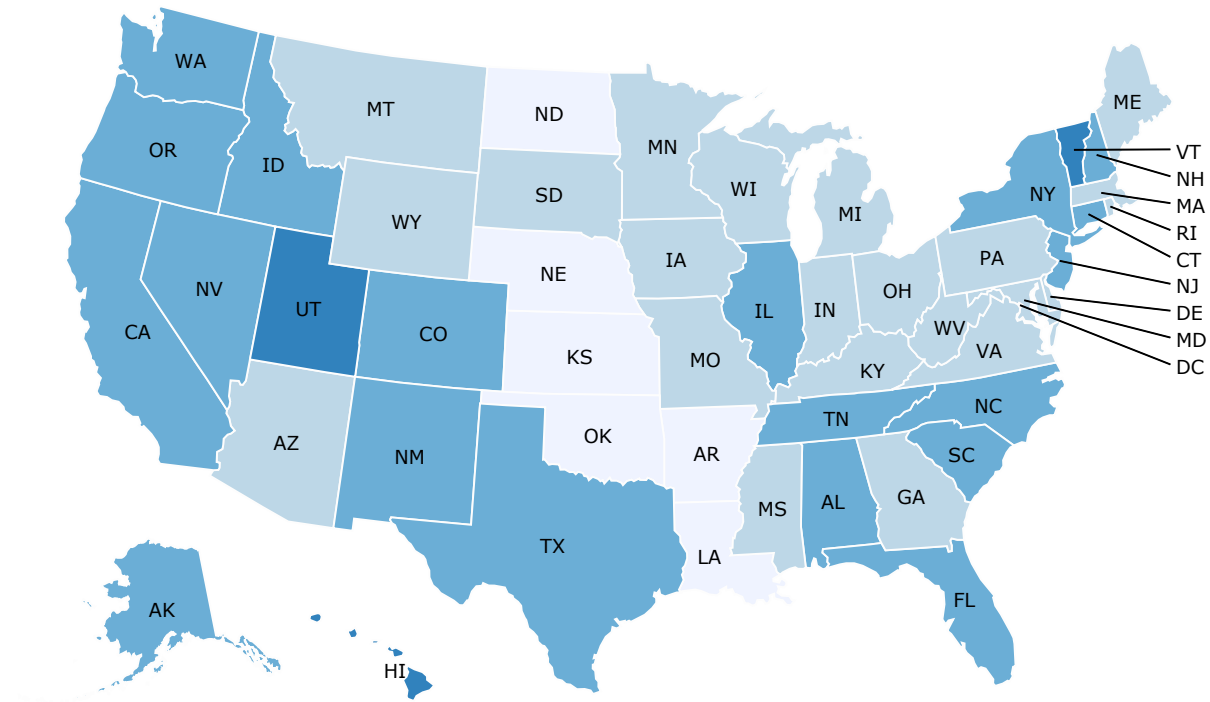
```
library(rMaps)
year_state_sum("dti") %>%
  mutate(year = as.numeric(orig_year)) %>%
  rMaps::ichoropleth(ave_dti ~ state, data = .,
    animate = "year",
    geographyConfig = list(popupTemplate = "#!function(geo, data) {
      return '<div class=\"hoverinfo\"><strong>' +
        data.state + '<br>' + 'Average DTI in ' + data.year + ': ' +
        data.ave_dti.toFixed(2) +
        '</strong></div>';}!#")) -> state_fico

state_fico$save("StateMapDTI.html", cdn = T)
```



# Plot

1999



# Machine Learning with **RxSpark**

# Predictive Models in the **RevoScaleR** package



## ETL

- Data import – Delimited, Fixed, SAS, SPSS, ODBC
- Variable creation & transformation
- Recode variables
- Factor variables
- Missing value handling
- Sort, Merge, Split
- Aggregate by category (means, sums)



## Descriptive Statistics

- Min / Max, Mean, Median (approx.)
- Quantiles (approx.)
- Standard Deviation
- Variance
- Correlation
- Covariance
- Sum of Squares (cross product matrix for set variables)
- Pairwise Cross tabs
- Risk Ratio & Odds Ratio
- Cross-Tabulation of Data (standard tables & long form)
- Marginal Summaries of Cross Tabulations



## Statistical Tests

- Chi Square Test
- Kendall Rank Correlation
- Fisher's Exact Test
- Student's t-Test



## Predictive Statistics

- Sum of Squares (cross product matrix for set variables)
- Multiple Linear Regression
- Generalized Linear Models (GLM) exponential family distributions: binomial, Gaussian, inverse Gaussian, Poisson, Tweedie. Standard link functions: cauchit, identity, log, logit, probit. User defined distributions & link functions.
- Covariance & Correlation Matrices
- Logistic Regression
- Predictions/scoring for models
- Residuals for all models



## Variable Selection

- Stepwise Regression



## Machine Learning

- Decision Trees
- Decision Forests
- Gradient Boosted Decision Trees
- Naïve Bayes



## Clustering

- K-Means



## Sampling

- Subsample (observations & variables)
- Random Sampling



## Simulation

- Simulation (e.g. Monte Carlo)
- Parallel Random Number Generation



## Custom Parallelization

- rxDataStep
- rxExec
- PEMA-R API



# Building Machine Learning Pipelines

## Functional Decompositions

```
estimate_model <- function(xdf_data = freddie[["train"]],  
                           form = make_form(xdf_data, depVar = "default_flag"),  
                           model = rxLogit, ...) {  
  
  rx_model <- model(form, data = xdf_data, ...)  
  
  return(rx_model)  
  
}
```



# Executing Our Models In Parallel

## Functionals

```
computeContext <- RxSpark(consoleOutput=TRUE,  
                           nameNode=myNameNode,  
                           port=myPort,  
                           executorCores=6,  
                           executorMem = "10g",  
                           executorOverheadMem = "5g",  
                           persistentRun = TRUE)  
  
rxSetComputeContext(computeContext)  
  
models <- list(rxLogit, rxDTree, rxDForest, rxBTrees)  
trained_models <- rxExec(estimate_model, model = rxElemArg(models))
```



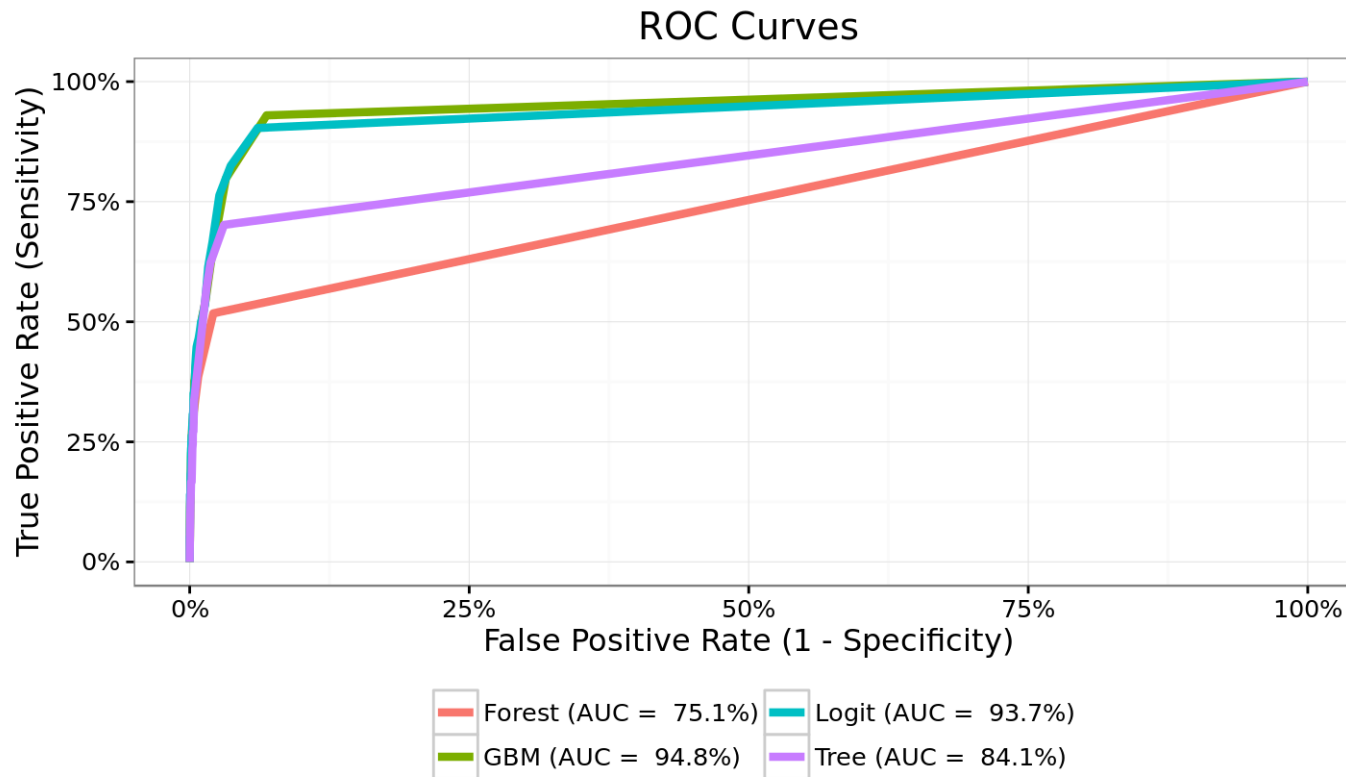
# Scoring the Models

```
default_pred <- function(model = default_model_gbm) {  
  
  scored_xdf <- rxPredict(model,  
                          mort_split$validate,  
                          "scored.xdf")  
  
  return(scored_xdf)  
}
```





# Model Results



What's Next

# Next Steps

- Sharing of Spark Contexts
  - currently, each of these APIs use their own backend to create Spark Sessions/Contexts, and cannot share objects across
- **RxSpark** will allow for using Spark DataFrames directly, and to share Spark sessions with other APIs
- Addition of new models and transformers in RevoScaleR



**Thank You!**

# Resources

- <https://github.com/akzaidi/R-cadence/tree/master/Spark>
- <https://bookdown.org/alizaidi/mrs-spark-ml/>
- <http://spark.rstudio.com/>

